

Introduction to Computer Graphics

6. GPU and Shaders

I-Chen Lin
National Chiao Tung University

Textbook: E. Angel, D. Shreiner Interactive Computer Graphics, 6th Ed., Pearson
Ref: D.D. Hearn, M. P. Baker, W. Carithers, Computer Graphics with OpenGL, 4th Ed., Pearson

The Development of Graphics Cards (consumer-level): Early 90's

- ▶ VGA cards in the early 90's
 - ▶ Just output designated “bitmap”.
 - ▶ Some with 2D acceleration, ex. “Bitblt”
 - ▶ Ex. S3
- ▶ Interactive 3D(or 2.5D) games relied on software rendering.
 - ▶ There were hardware graphics pipelines on workstations, e.g. SGI.

The Development of Graphics Cards (consumer-level): Late 90's

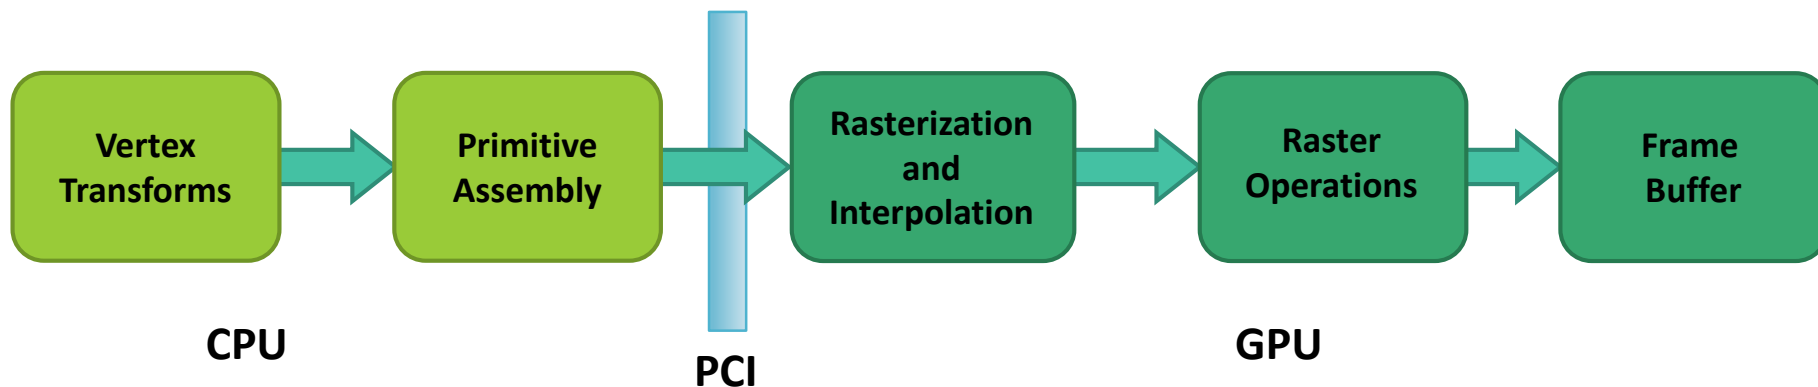
- ▶ 3D accelerators (90's)
 - ▶ Fixed-function pipelines.
 - ▶ E.g. S3, Voodoo, Nvidia, ATI, 3D Labs....
 - ▶ Some of them had to work with a standard VGA card.

3Dfx Voodoo (1996)

- ▶ One of the first true 3D game cards
- ▶ Worked by supplementing a standard 2D video card.
- ▶ Did not do vertex transformations (they were evaluated in the CPU)
- ▶ Did texture mapping, z-buffering.



en.wikipedia.org/wiki/3dfx_Interactive



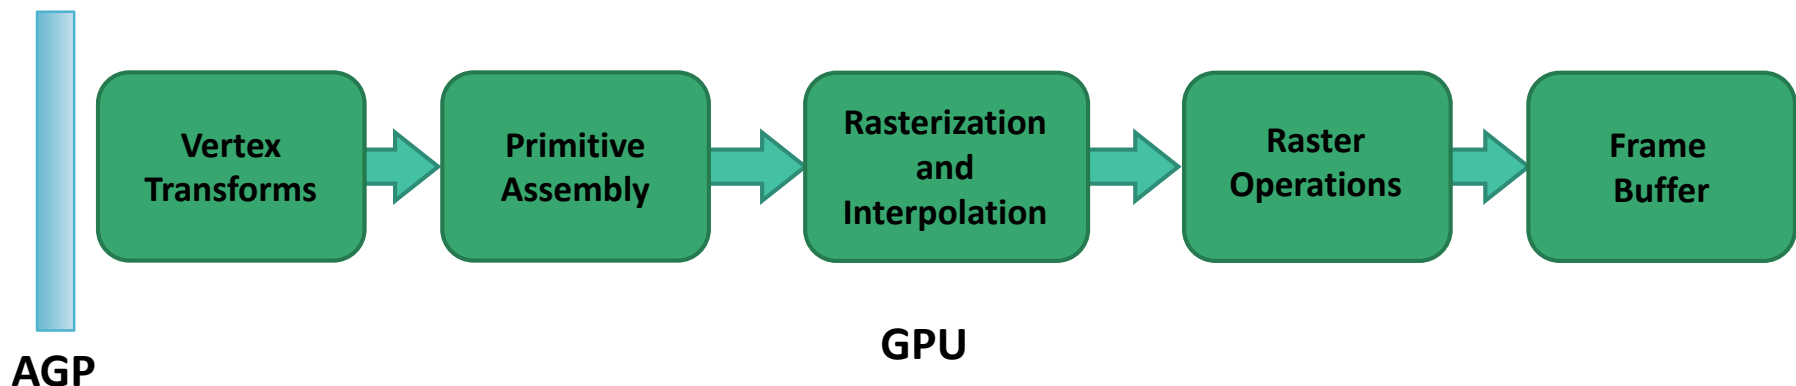
Modified from S. Venkatasubramanian and J. Kider, "Evolution of the Programmable Graphics Pipeline"

GeForce/Radeon 7500 (1998)

- ▶ Main innovation: shifting the transformation and lighting calculations to the GPU
- ▶ Allowed multi-texturing: giving bump maps, light maps, and others.
- ▶ Faster AGP bus instead of PCI



en.wikipedia.org/wiki/GeForce_256



The Development of Graphics Cards (consumer-level): after 2001

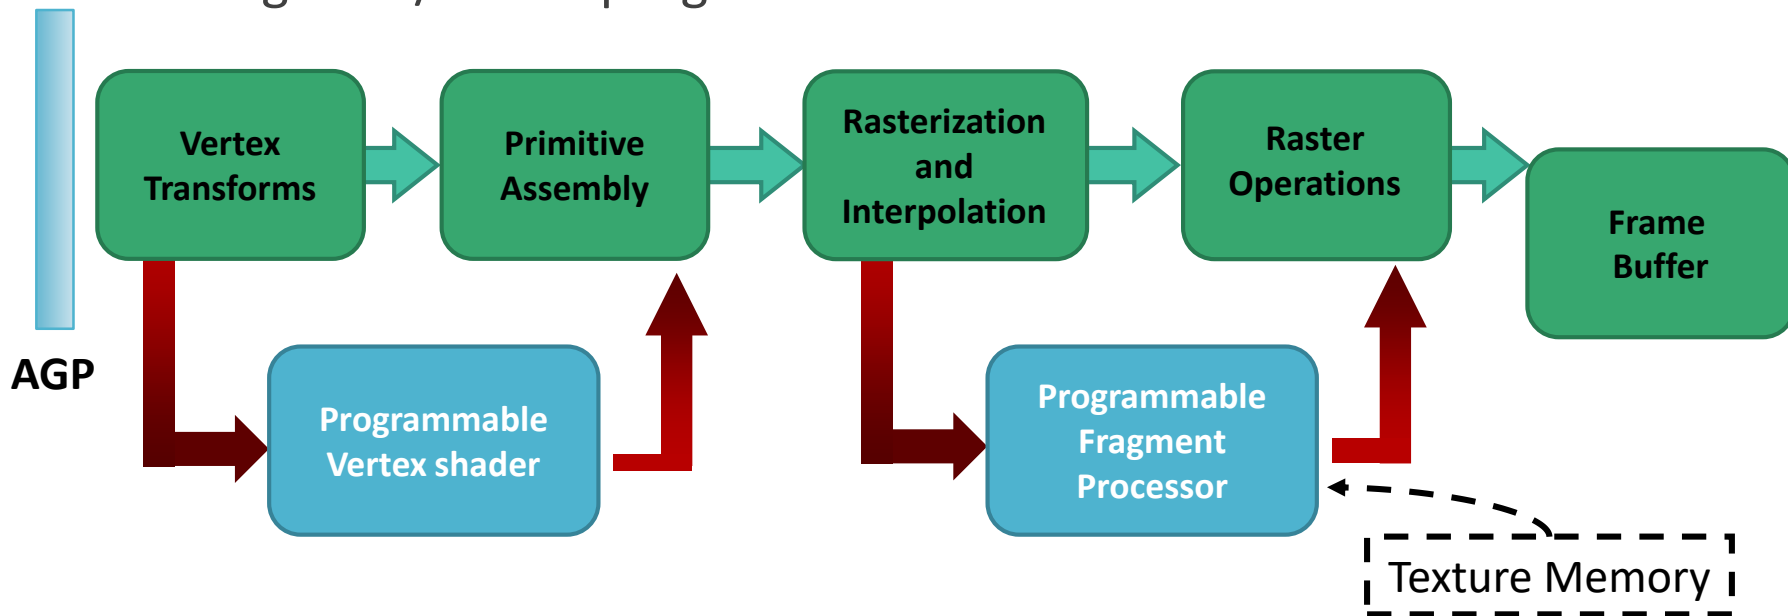
- ▶ Programmable pipelines on GPU
- ▶ GeForce3/Radeon 8500(2001)
 - ▶ Programmable vertex computations: up to 128 instructions
 - ▶ Limited programmable fragment computations: 8-16 instructions



https://en.wikipedia.org/wiki/GeForce_3_series

The Development of Graphics Cards (consumer-level): after 2001 (cont.)

- ▶ Radeon 9700/GeForce FX (2002)
 - ▶ the first generation of fully-programmable graphics cards
 - ▶ Different versions have different resource limits on fragment/vertex programs



Modified from S. Venkatasubramanian and J. Kider, "Evolution of the Programmable Graphics Pipeline"

Evaluation of Graphics Pipeline

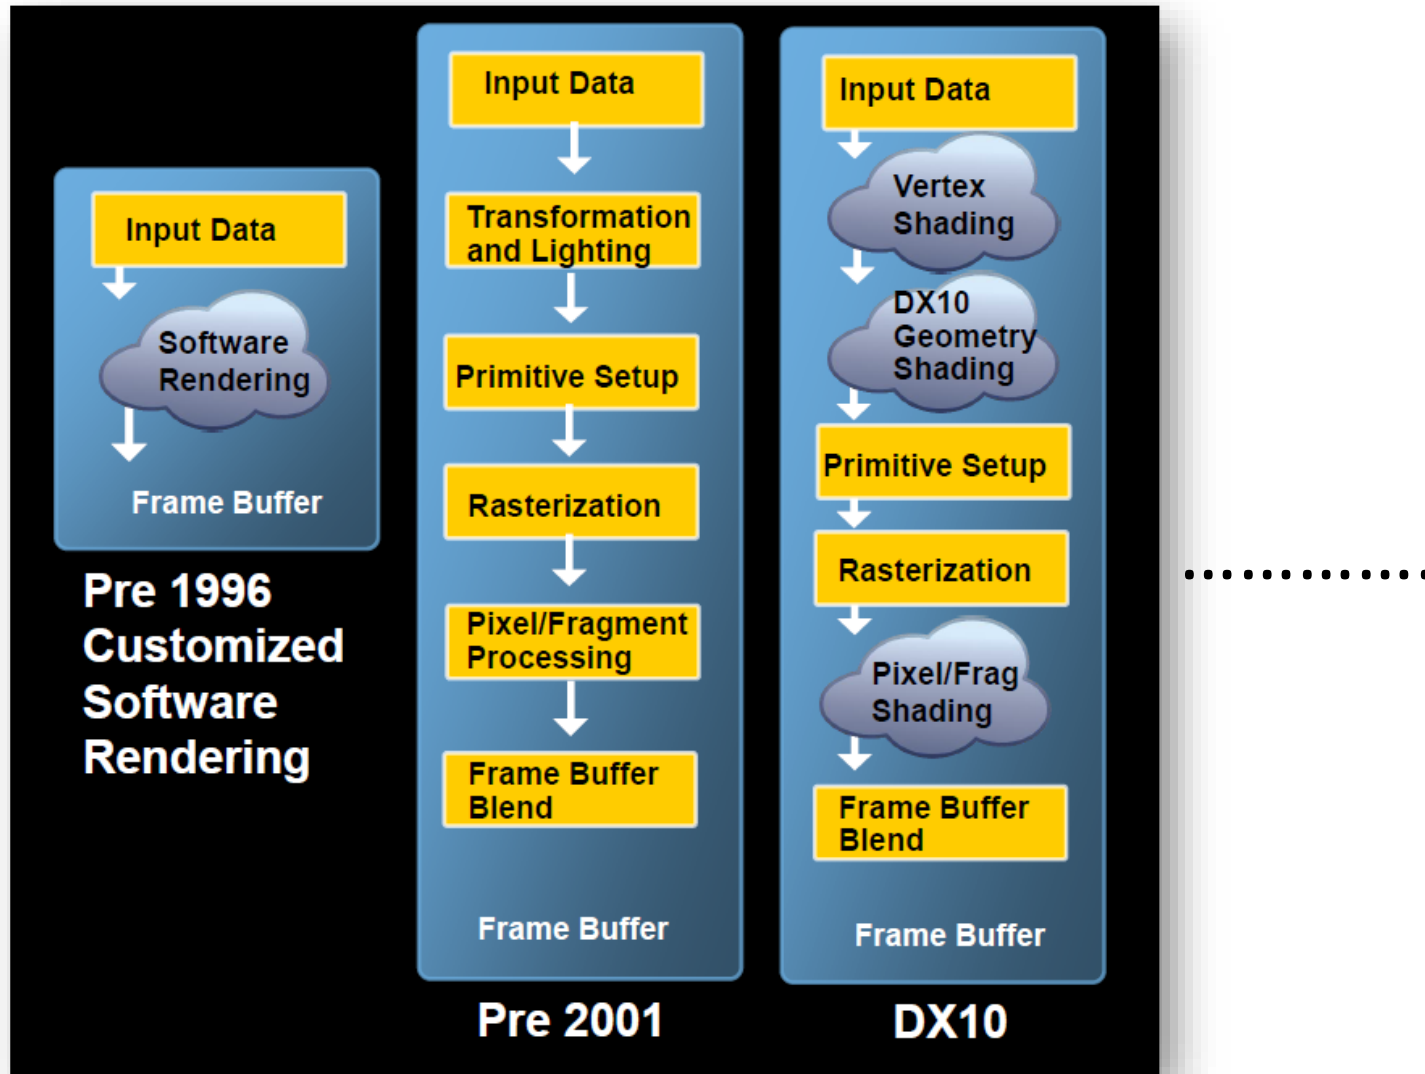


Figure from: M. Houston, “Beyond Programmable Shading Retrospective” slides

GPU & Shaders : the new age of real-time graphics

- ▶ Programmable pipelines.
- ▶ Supported by high-end commodity cards
 - ▶ NVIDIA, AMD/ATI, etc.



en.wikipedia.org/wiki/GeForce_10_series



www.amd.com/zh-hant/products/graphics/radeon-rx-570

Why is It So Remarkable?

- ▶ We can do lots of cool stuff in real-time, without overworking the CPU.
 - ▶ Phong Shading
 - ▶ Bump Mapping
 - ▶ Particle Systems
 - ▶ Animation
 - ▶
- ▶ Beyond real-time graphics: GP-GPU, e.g. CUDA, OpenCL (Open Computing Language)
 - ▶ Scientific Data Processing
 - ▶ Computer vision
 - ▶ Deep learning
 - ▶

Programmable Components

- ▶ **Shader:** programmable processors.
 - ▶ Replacing fixed-function vertex and fragment processing, and so forth.
- ▶ Types of shaders:
 - ▶ **Vertex shaders**
 - ▶ Dealing with per-vertex functions.
 - ▶ We can control the lighting and position of each vertex.
 - ▶ **Fragment shaders**
 - ▶ Dealing with per-pixel functions.
 - ▶ We can control the color of each pixel by user-defined programs.
 - ▶ **Geometry shaders** (DirectX 10, SM 4+)
 - ▶ New shaders (hull, domain) in DirectX11, SM5

Programmable Components (cont.)

- ▶ Software Support

- ▶ Direct X 8 , 9, 10, 11, 12, ...

- ▶ OpenGL Extensions

- ▶ OpenGL Shading Language (GLSL)

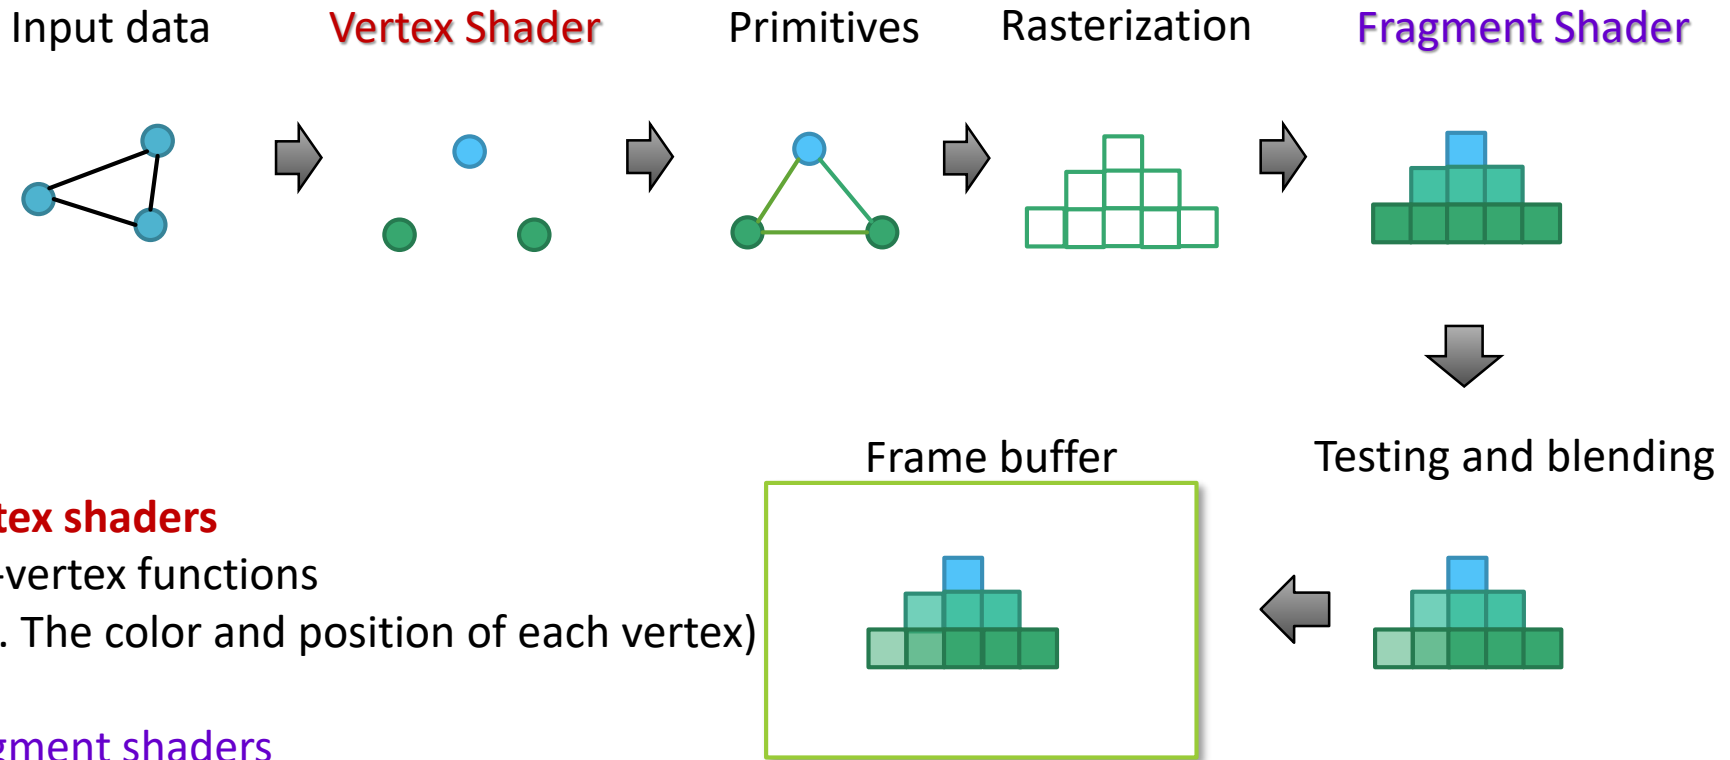
- ▶ OpenGL for Embedded Systems (OpenGL ES)

- ▶ *Cg (C for Graphics)*

- ▶ Metal Shading Language (by Apple)

- ▶

Essential GLSL pipeline (Vert.+Frag. Shaders)



Vertex shaders

per-vertex functions
(E.g. The color and position of each vertex)

Fragment shaders

per-fragment (pixel) function.
(E.g. The color of each fragment)

What about GLSL programs?

- ▶ Besides your main program (e.g. main.cpp), there are additional shader codes.
- ▶ These code can be character strings in your .cpp, but we usually put them in separate files (e.g. ooo.vs or ooo.vert, xxx.fs or xxx.frag).
- ▶ In a GLSL program, you can use multiple (different) shader codes to demonstrate different illumination algorithms for objects or regions. 一個shader不是搭配一個function · 可用不同shader對應同一物體

Vert buffer passed
from main.cpp

*不能用glBegin、glEnd找點：效率較差

A simple example of shader codes

宣告版本：330核心版

Vertex shader code

```
#version 330 core
layout (location = 0) in vec3 aPos; // the position variable has attribute position 0
```

```
out vec4 vertexColor; // specify a color output to the fragment shader output四維顏色
```

```
void main()
{
    gl_Position = vec4(aPos, 1.0); // see how we directly give a vec3 to vec4's constructor
    vertexColor = vec4(0.5, 0.0, 0.0, 1.0); // set the output variable to a dark-red color
}
```

這個一定要放：收進的position · output會變甚麼

收集到的顏色直接往下塞

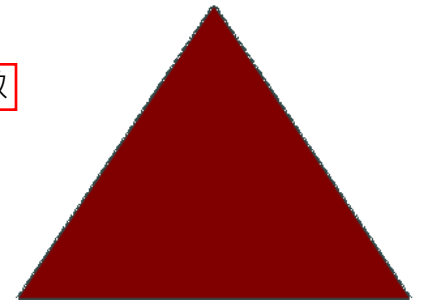
Fragment shader code

```
#version 330 core
out vec4 FragColor; // 一定要宣告
```

```
in vec4 vertexColor; // the input variable from the vertex shader in/out要一致
```

```
void main()
{
    FragColor = vertexColor;
}
```

```
float vertices[] = {
VAO  0.5f, -0.5f, 0.0f,
VBO  -0.5f, -0.5f, 0.0f,
      0.0f, 0.5f, 0.0f
}; // input是VBO ·
    裡面會存個array
```



Note: **gl_FragColor** is deprecated The example is modified from samples in learnopengl.com

Another simple example

基本上照抄

Codes in main.cpp

```
// For the vertex shader
```

```
int vertexShader = glCreateShader(GL_VERTEX_SHADER);
```

```
glShaderSource(vertexShader, 1, &vertexShaderSource, NULL);
```

從某個character string而來

```
glCompileShader(vertexShader); => 在執行前compile
```

.....

```
// For the fragment shader
```

```
int fragmentShader = glCreateShader(GL_FRAGMENT_SHADER);
```

```
glShaderSource(fragmentShader, 1, &fragmentShaderSource, NULL);
```

```
glCompileShader(fragmentShader);
```

一般從file讀進來(或character string而來)

.....

```
// link the above shaders
```

```
int shaderProgram = glCreateProgram();
```

```
glAttachShader(shaderProgram, vertexShader);
```

連接到主程式

```
glAttachShader(shaderProgram, fragmentShader);
```

```
glLinkProgram(shaderProgram);
```

.....

The example is modified from samples in learnopengl.com

Another simple example (cont.)

```
float vertices[] = {  
    // positions // colors  
    0.5f, -0.5f, 0.0f, 1.0f, 0.0f, 0.0f, // bottom right  
    -0.5f, -0.5f, 0.0f, 0.0f, 1.0f, 0.0f, // bottom left  
    0.0f, 0.5f, 0.0f, 0.0f, 0.0f, 1.0f // top  
};
```

```
unsigned int VBO, VAO;
```

```
glGenVertexArrays(1, &VAO);
```

```
glGenBuffers(1, &VBO);
```

```
glBindVertexArray(VAO);
```

```
glBindBuffer(GL_ARRAY_BUFFER, VBO);
```

```
glBufferData(GL_ARRAY_BUFFER, sizeof(vertices), vertices, GL_STATIC_DRAW);
```

```
// position attribute position/element/float//每次往後跳，要經過6float(offset要自己算)/起始點從0開始
```

```
glVertexAttribPointer(0, 3, GL_FLOAT, GL_FALSE, 6 * sizeof(float), (void*)0);
```

```
glEnableVertexAttribArray(0);
```

*放array和offset要小心

```
// color attribute position/element/float//每次往後跳，要經過6float(offset要自己算)/起始點從第三格開始
```

```
glVertexAttribPointer(1, 3, GL_FLOAT, GL_FALSE, 6 * sizeof(float), (void*)(3 * sizeof(float)));
```

```
glEnableVertexAttribArray(1);
```

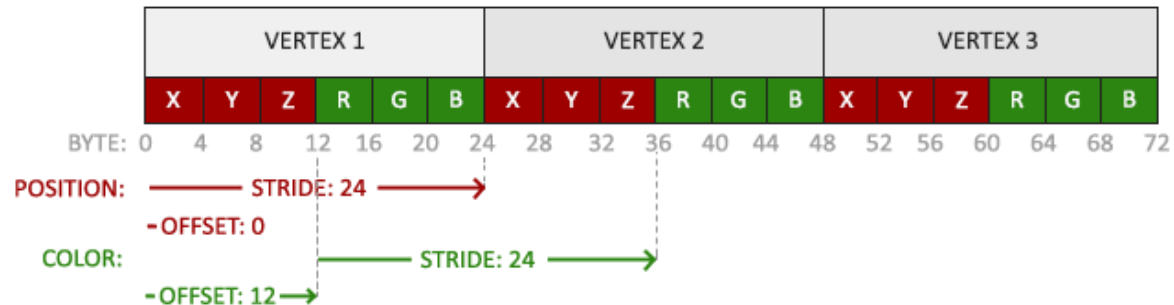
.....

```
glUseProgram(shaderProgram); => 選擇要使用哪個shader
```

Codes in main.cpp

VAO : vertex array buffer

VBO : 實際在放的buffer data(EX : vertex shader)



Another simple example (cont.)

Vertex shader code

```
#version 330 core
```

```
layout (location = 0) in vec3 aPos; // the position variable at position 0 [0.5, -0.5, 0]
layout (location = 1) in vec3 aColor; // the color variable at position 1 [0, 1, 0]
```

每個點會獨立收到一個不同的資料 attributePosition/attributeColor

```
out vec3 ourColor; // output a color to the fragment shader
```

```
void main()
```

```
{ gl_Position = vec4(aPos, 1.0);
  ourColor = aColor; // set ourColor to the input vertex color
}
```

out attribute color會做內插
(vs. phong shading : 用out normal做內插)

Fragment shader code

```
#version 330 core
```

```
out vec4 FragColor;
```

```
in vec3 ourColor; [內插完的結果]
```

```
void main()
```

```
{ FragColor = vec4(ourColor, 1.0); 直接塞color轉成四維
}
```

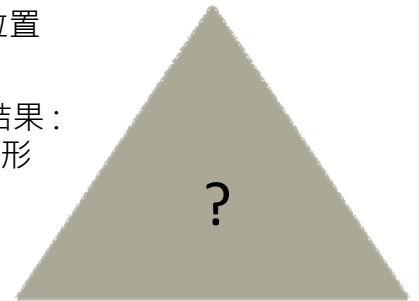
Vert buffer passed from main.cpp

```
float vertices[] = {
    // positions    // colors
    0.5f, -0.5f, 0.0f, 1.0f, 0.0f, 0.0f,
    -0.5f, -0.5f, 0.0f, 0.0f, 1.0f, 0.0f,
    0.0f, 0.5f, 0.0f, 0.0f, 0.0f, 1.0f
};
```

每個fragment都會平行執行fragment shader ·
自己只知道自己的資料(點) · 只能動顏色、不能動位置

(0, 0, 1) => blue

Rasterize結果：
漸層式三角形
(三色漸變)
(三點RGB)



(0, 1, 0) => green

(1, 0, 0) => red

An example with **matrix passing**

```
float vertices[] = {  
    // positions    // colors  
    0.5f, -0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // Mid right  
    -0.5f, -0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // Mid left  
    0.0f, 0.5f, -1.5f, 0.5f, 0.0f, 0.0f, // top  
    -0.5f, -0.5f, -1.5f, 0.0f, 0.5f, 0.0f, // Mid left  
    0.5f, -0.5f, -1.5f, 0.0f, 0.5f, 0.0f, // Mid right  
    0.0f, -1.5f, -1.5f, 0.0f, 0.5f, 0.0f // bottom  
};
```

Codes in main.cpp

Data passing through
uniform variables *glm::mat4

```
.....  
glClearColor(0.0f, 0.0f, 0.0f, 1.0f);  
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
```

```
glm::mat4 projection_matrix = glm::perspective(glm::radians(60.0f), (float)SCR_WIDTH / (float)SCR_HEIGHT,  
0.1f, 10.0f);
```

```
glUniformMatrix4fv(glGetUniformLocation(shaderProgram, "projection"), 1, GL_FALSE,  
glm::value_ptr(projection_matrix));
```

link到shader code的變數名稱

另一種寫法：&projection_matrix(0)(0);

```
glm::mat4 model_matrix = glm::mat4(1.0f); // make sure to initialize matrix to identity matrix first
```

```
model_matrix = glm::rotate(model_matrix, glm::radians(15.0f), glm::vec3(0.0f, 0.0f, 1.0f));
```

```
glUniformMatrix4fv(glGetUniformLocation(shaderProgram, "model"), 1, GL_FALSE, &model_matrix[0][0]);
```

↑
原來的matrix再乘matrix(model沿z軸逆時針旋轉15度)

送到shader

```
glDrawArrays(GL_TRIANGLES, 0, 6); 6 : 6個頂點=>兩個三角形(頂點重複)
```

An example with matrix passing (cont.)

```
#version 330 core
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aColor;
out vec3 ourColor;
uniform mat4 model;
uniform mat4 projection;
```

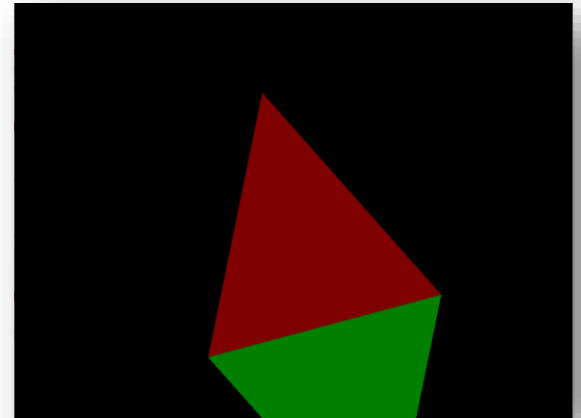
Vertex shader code

先傳誰進來沒差，乘的順序對就好

```
void main()
{
    gl_Position = projection * model * vec4(aPos, 1.0);
    ourColor = aColor;
}
```

Fragment shader code

```
#version 330 core
out vec4 FragColor;
in vec3 ourColor;
void main()
{
    FragColor = vec4(ourColor, 1.0f);
}
```



Vertex Shaders

- ▶ Per-vertex calculations performed here
 - ▶ Without knowledge about other vertices (parallelism)
- ▶ Your program take responsibility for:
 - ▶ Vertex transformation
 - ▶ Normal transformation
 - ▶ (Per-Vertex) Lighting
 - ▶ Color material application and color clamping
 - ▶ Texture coordinate generation

堆疊

Vertex Shader Applications

- ▶ We can control movement with uniform variables and vertex attributes
 - ▶ Time
 - ▶ Velocity
 - ▶ Gravity
- ▶ Moving vertices
 - ▶ Morphing
 - ▶ Wave motion
 - ▶
- ▶ Lighting
 - ▶ More realistic models
 - ▶ Cartoon shaders

Applications: Wave Motion Vertex Shader

Uniform: passing parameters to vertex and fragment shaders.

```
uniform float time;  
uniform float xs, zs;
```

```
void main()
```

```
{vec3 object_pos;
```

```
float s;
```

```
s = 1.0 + 0.1*sin(xs*time)*sin(zs*time);
```

 用layout傳：每個點都要拿到不同的sin的點

```
object_pos = aPos;
```

```
object_pos.y = s* aPos.y;
```

 現在禁用(因為效率造成影響)

```
gl_Position = gl_ModelViewProjectionMatrix* object_pos;
```

```
}
```

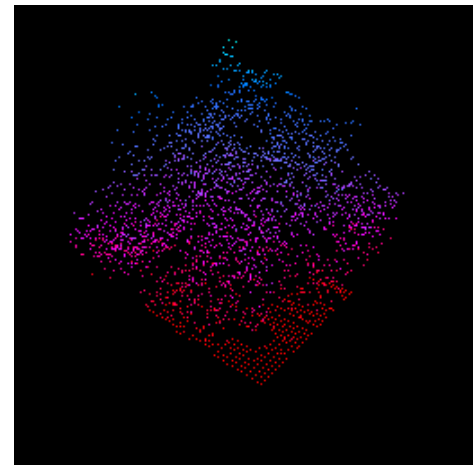
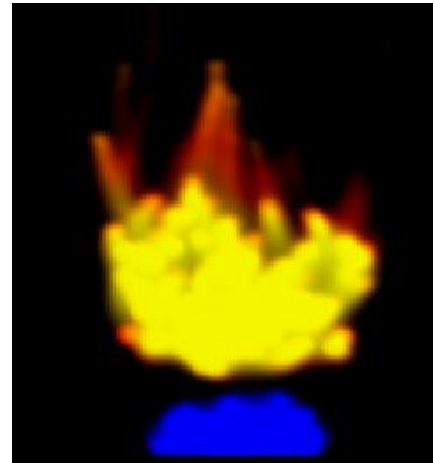
改變時間t

Note: Several `gl_` predefined variables are deprecated in the newer version. Use uniform variables instead.

Applications: Particle Systems

Uniform: passing parameters to vertex and fragment shaders.

```
.....  
uniform vec3 vel;  
uniform float g, m, t;  
  
void main() 只變更時間t  
{  
    vec3 object_pos;  
    object_pos.x = aPos.x + vel.x*t;  
    object_pos.y = aPos.y + vel.y*t + g/(2.0*m)*t*t;  
    object_pos.z = aPos.z + vel.z*t;  
    gl_Position = gl_ModelViewProjectionMatrix*  
    vec4(object_pos,1);  
}
```



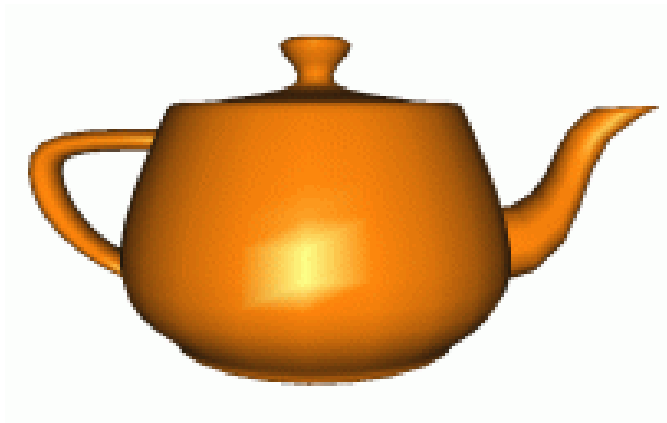
Note: Several gl_ predefined variables are deprecated in the newer version. Use uniform variables instead.

Fragment Shaders

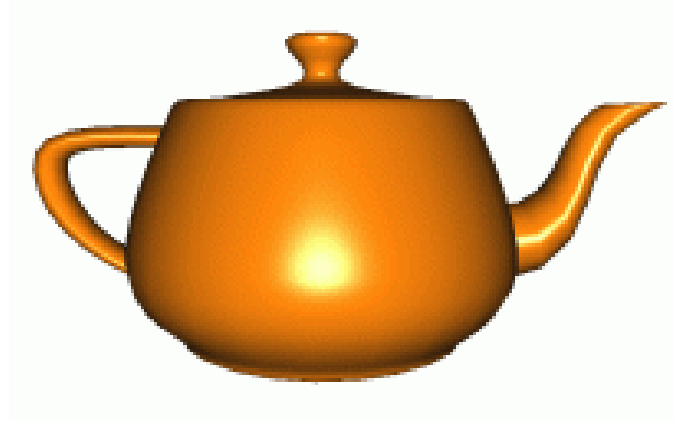
- ▶ What is a fragment? 可能很多點疊在一起，但最終只會在螢幕上顯示一個點
 - ▶ Cg Tutorial says: “You can think of a fragment as a ‘potential pixel’”
- ▶ Perform per-pixel calculations
 - ▶ Without knowledge about other fragments (parallelism)
- ▶ Your program’s responsibilities:
 - ▶ Operations on interpolated values
 - ▶ Texture access and application
 - ▶ Other functions: fog, color lookup, etc.

Fragment Shader Applications

(Per-pixel) Phong shading 每個三角形裡都要算顏色



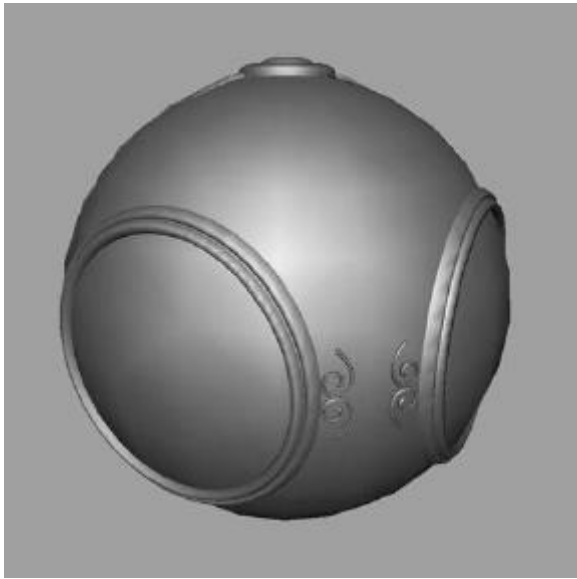
Per-vertex lighting



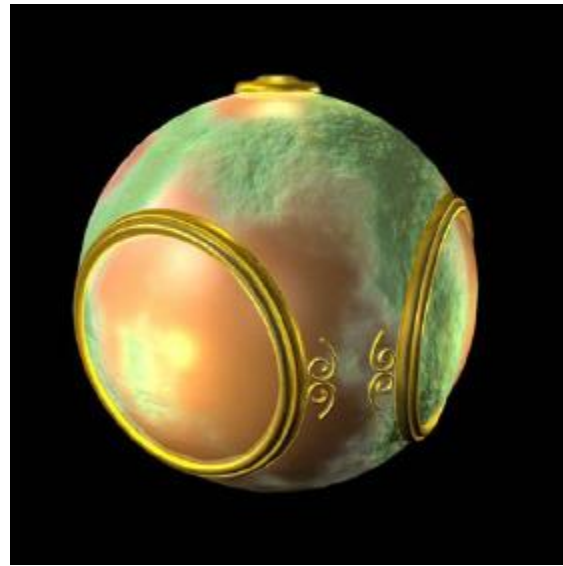
Per-fragment lighting

Figures from <http://www.lighthouse3d.com/opengl/glsl/>

Fragment Shader Applications



smooth shading

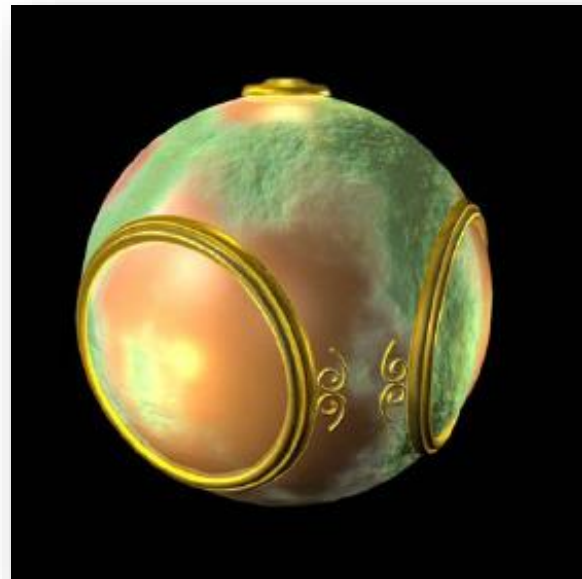
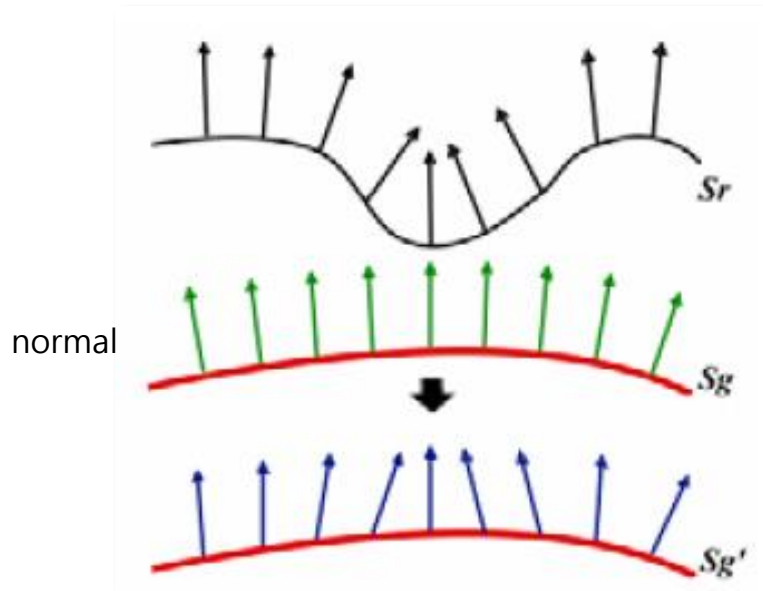


bump mapping

Bump Mapping

擾動

- ▶ Perturb normal for each fragment
- ▶ Store perturbation as textures



節省texture memory、顏色階段化

Toon Shading

卡通渲染

- The vertex shader then becomes:

.....

```
out vec3 vnormal; 先做phong shading壓成三種顏色
```

```
void main() {
```

```
    vnormal = gl_NormalMatrix * gl_Normal;
```

```
    gl_Position = ftransform(); }
```

要往下傳norma

=>改用projection * view * model * apos

- The fragment shader becomes

.....

```
in vec3 vnormal;
```

```
out vec4 FragColor;
```

```
void main() {
```

```
    float intensity; vec4 color;
```

```
    vec3 n = normalize(vnormal);
```

```
    intensity = dot(vec3(gl_LightSource[0].position),n); 也被禁用=>改light * normal
```

```
    if (intensity > 0.95) color = vec4(1.0,0.5,0.5,1.0);
```

```
    else if (intensity > 0.5) color = vec4(0.6,0.3,0.3,1.0);
```

```
    else if (intensity > 0.25) color = vec4(0.4,0.2,0.2,1.0);
```

```
    else color = vec4(0.2,0.1,0.1,1.0);
```

```
    FragColor = color; }
```

現在被禁用惹

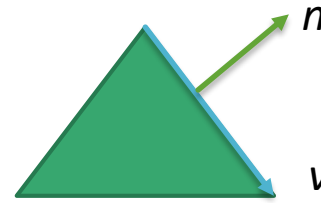
ftransform(): result from the GL fixed-function transformation pipeline

*Note: **varying**, communicating between vertex and fragment. Use **in out** variables in newer versions.*



Example from <http://www.lighthouse3d.com/opengl/gls/>

gl_NormalMatrix



$$n \cdot v = 0$$

*自己想旋轉 : normal * rotation

► Can we directly apply the modelview matrix M to a normal vector ?

► Problem: If the upper-left 3×3 submatrix M_s is **not orthogonal**, $n' = M_s n$ is not perpendicular to $v' = M_s v$

*Our goal is to find a matrix N for $n'' \cdot v' = 0$, where $n'' = Nn$. Nn未知

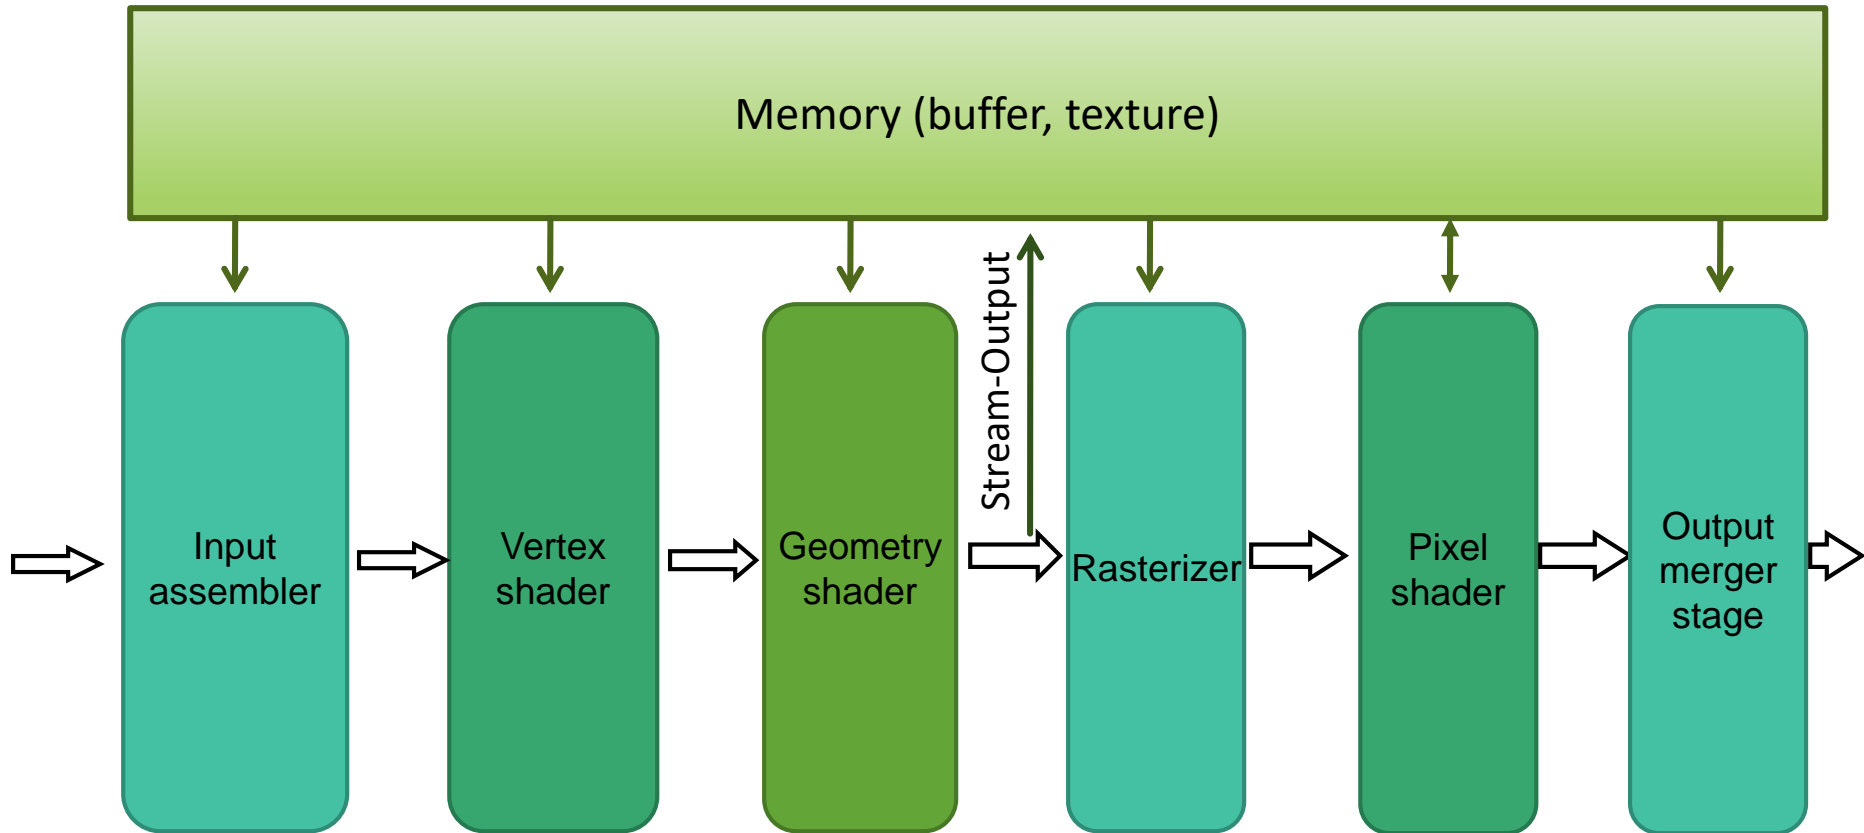
*(Nn) dot (Msv) = 0 = $(n^t)(N^t)Msv$

*It is reasonable to choose $(N^t)Ms = I$, since $(n^t)v = 0$

互為inverse

$$N = (M_s)^{-1^t}$$

With the Geometry Shader



Direct3D 10 pipeline stage from MSDN of Microsoft

D3D 10 Pipeline

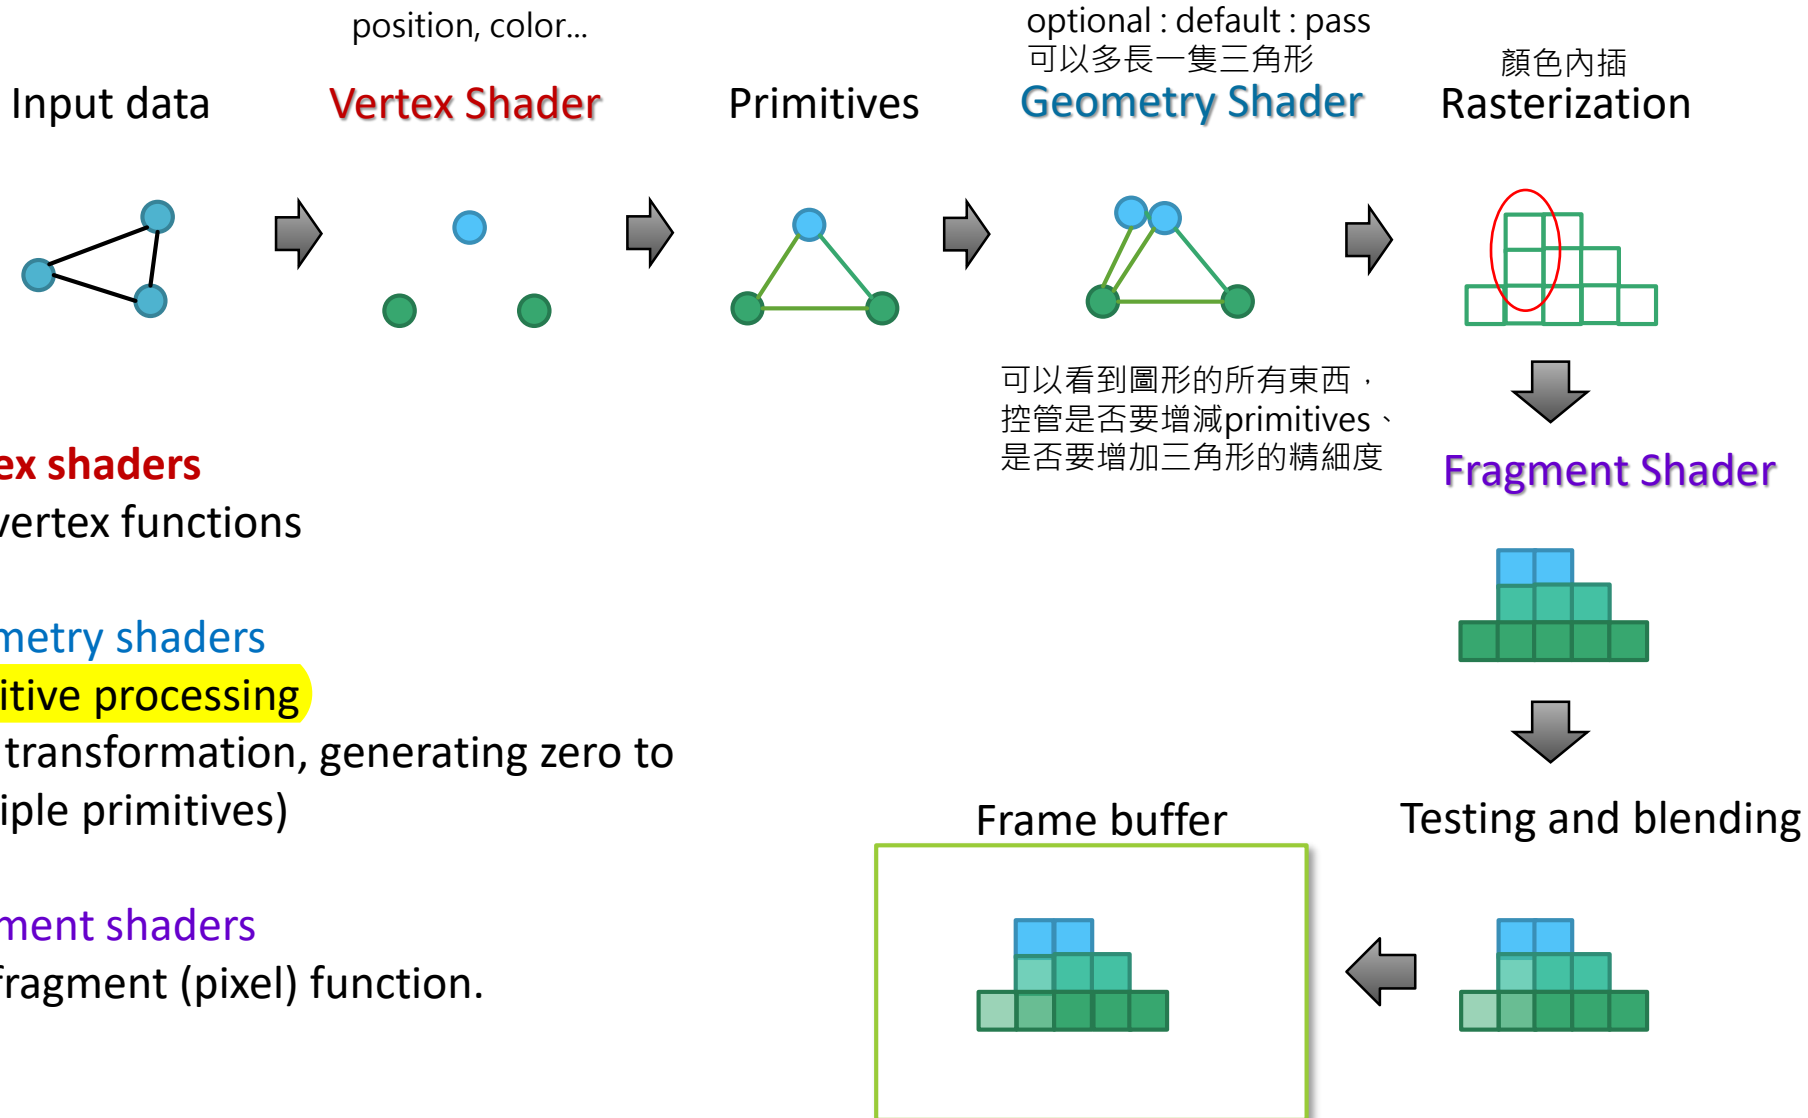
- ▶ **Input assembler:** supplies data (triangles, lines and points) to the pipeline.
- ▶ **Vertex shader:** processes vertices, such as transformations, skinning, and lighting.
- ▶ **Geometry shader:** processes entire primitives.
 - ▶ 3 vertices: a triangle, 2 vertices: a line, or 1 vertex: a point.
放大
 - ▶ The Geometry shader supports limited geometry amplification and de-amplification. (discard the primitive, or emit one or more new primitives)
一次可以看整個三角形或邊，也可以多加點
 - ▶ E.g. Subdivision, point -> billboard, silhouette edge -> fur, etc.
subdivision: 一個三角形拆成三個三角形 邊緣長毛
- ▶ **Stream-output stage:**
 - ▶ Data can be streamed out and/or passed into the rasterizer. Data streamed out to memory can be recirculated back into the pipeline as input data or read-back from the CPU.

D3D 10 Pipeline (cont.)

position中間做內插

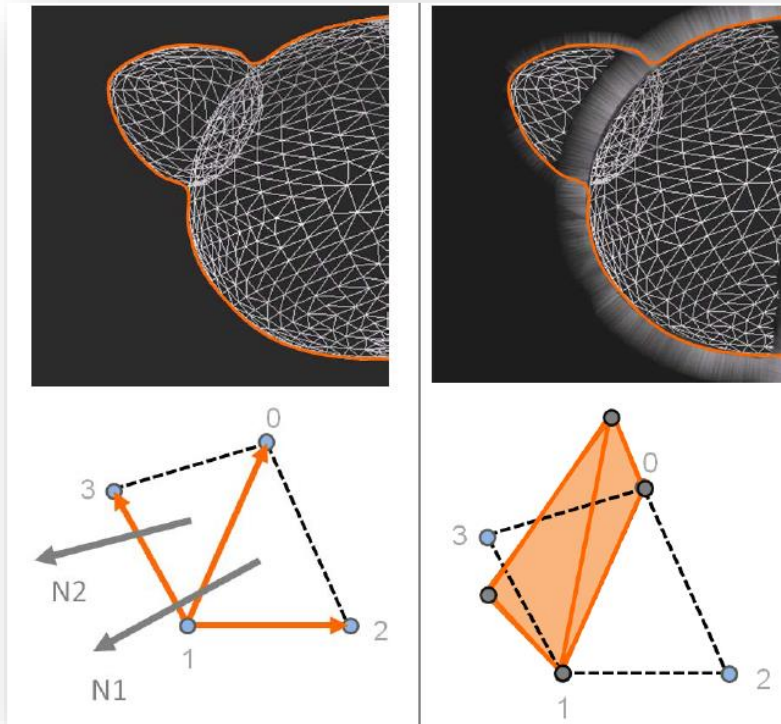
- ▶ **Rasterizer:** clips primitives, prepares primitives for the pixel shader and determines how to invoke pixel shaders.
- ▶ **Pixel shader:** receives interpolated data for a primitive and generates per-pixel data, such as color.
- ▶ **Output-merger stage:** 三角形做混合或把不用的東西砍掉
 - ▶ combines various types of output data (pixel shader values, depth and stencil information) with the contents of the render target and depth/stencil buffers to generate the final pipeline result.

GLSL pipeline (Vert.+Geo.+Frag. Shaders)

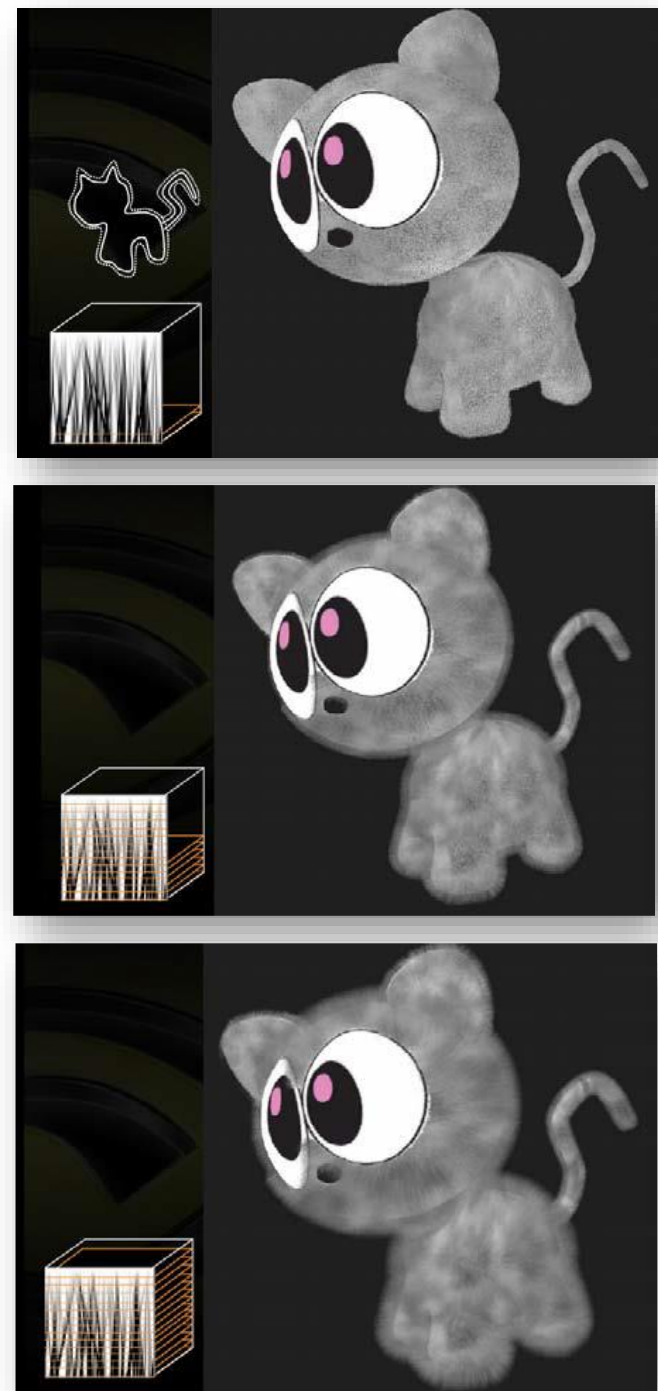


D3D 10 Pipeline (cont.)

細毛



Figures from NVIDIA DirectX10 SDK Doc:
Fur (using Shells and Fins)



vertex=>geometry=>fragment

Previous example using Geometry Shader

Vertex shader code

Geometry shader code

```
#version 330 core
layout (location = 0) in vec3 aPos;
layout (location = 1) in vec3 aColor;

out VS_OUT{
    vec3 vsColor;
} vs_out;

uniform mat4 model;
uniform mat4 projection;
```

兩個三角形一邊(兩點)重疊

```
void main()
{ gl_Position = projection * model * vec4(aPos, 1.0);
  vs_out.vsColor = aColor;
}
```

Fragment shader code

```
#version 330 core
out vec4 FragColor;
in vec3 ourColor;

void main()
{ FragColor = vec4(ourColor, 1.0f);
}
```

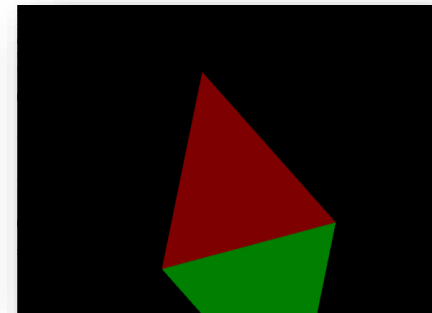
```
#version 330 core
layout (triangles) in;
layout (triangle_strip, max_vertices = 3) out;

in VS_OUT{
    vec3 vsColor;
} gs_in[];
out vec3 ourColor;
```

得到兩個分開的三角形

```
void main()
{
    for(int i=0; i<3; i++)
    { gl_Position = gs_in[i].gl_Position;
      ourColor = gs_in[i].vsColor;
      EmitVertex();
    }
    EndPrimitive();
}
```

把資料收進來，
normal可以在這裡算



input 3 vertex, 1 triangle => input 4 vertex, 2 triangle

Adding additional triangles with GS

Geometry shader code

可以省bandwidth

```
#version 330 core
layout (triangles) in;
layout (triangle_strip, max_vertices = 4) out;
in VS_OUT{
    vec3 vsColor;
} gs_in[];
out vec3 ourColor;
```

新增一個點

Note:

Triangle strip: v0, v1, v2, v3
⇒ Triangle 1 (v0, v1, v2)
⇒ Triangle 2 (v1, v2, v3)

```
void main() 把資料抓進來(原始三個點)
{ for(int i=0; i<3; i++)
{ gl_Position = gl_in[i].gl_Position;
  ourColor = gs_in[i].vsColor;
  EmitVertex(); 把資料吐出去
} normalized space vertex
gl_Position = gl_in[0].gl_Position + vec4(0.2f, 0.2f, -0.2f, 0.0f);
ourColor = vec3(0.0f, 0.0f, 0.8f);
EmitVertex();
EndPrimitive();
}
```

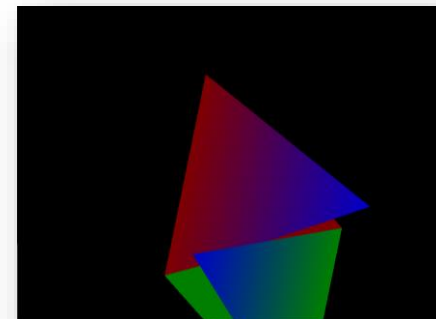
projection * model matrix => static constant
在原始空間操作完(EX: 加一個點) ·
再放Geometry shader

uniform可以送給所有shader

For each triangle, add one additional triangle.

在這個地方多塞一個點

多長了一個斜斜的三角形



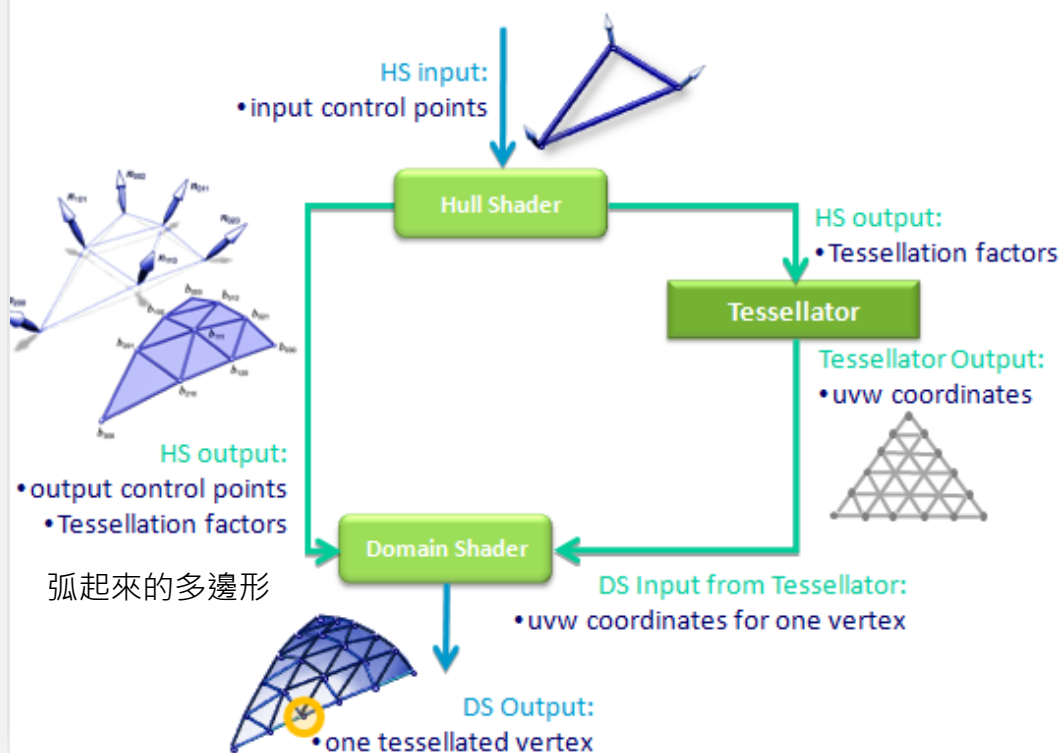
D3D 11 Pipeline

- ▶ In D3D10, the Geometry shader may subdivide the surfaces by multiple passes.

把東西細化・把三角形割碎

- ▶ D3D11 improves the **tessellation** ability by three new stages: hull shader, tessellator, domain shader.
- ▶ The tessellated patches can still be applied to geometry shaders. E.g. point -> billboard, silhouette edge -> fur, etc.

Tessellation Pipeline



input assembler
vertex shader
tessellation
geometry shader
rasterization
fragment shader
color blending
framebuffer

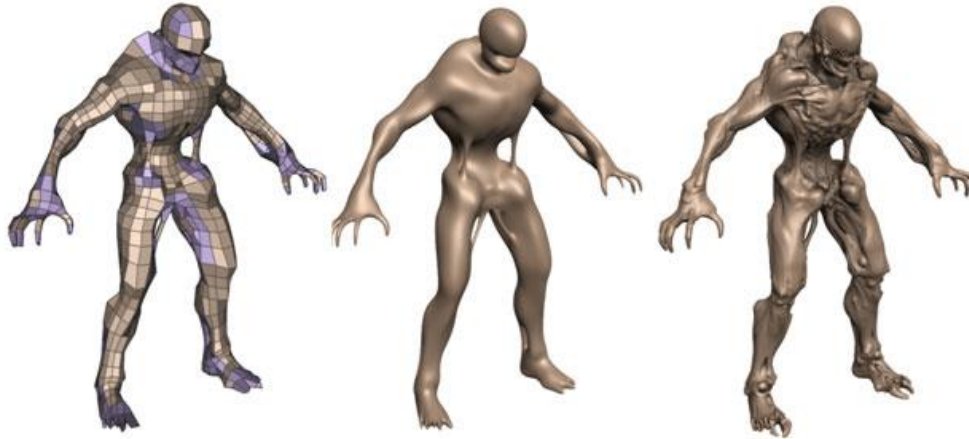
Figure from:
developer.download.nvidia.com/presentations/2009/GDC/GDC09_D3D11Tessellation.pdf

Figure from: vulkan-tutorial.com/Drawing_a_triangle/Graphics_pipeline_basics/Introduction

D3D 11 Tessellation



Model refinement



Tessellation with displacement mapping

End of Chapter 6